

# Homework 2, STAT 5241

Zongyi Liu

Mar 1, 2025

## 1 Questions

### 1.1 Question 1

To address the challenge of predicting autism in a European population using a model trained on US data, the team must carefully design their machine learning process to account for the significant differences in key features such as ethnicity, demographics, and environmental factors. The first step involves thorough data preprocessing to align and normalize features across the two datasets, ensuring consistency and handling missing values through imputation. Feature engineering should focus on creating universal patterns rather than region-specific ones. Given the distribution shift between the training and testing datasets, the team should employ domain adaptation techniques such as transfer learning, domain-adversarial neural networks (DANN), or covariate shift correction to improve generalization. Robust models like random forests, gradient boosting machines (e.g., XGBoost, LightGBM), or regularized linear models are recommended due to their ability to handle mixed data types and resist overfitting to region-specific patterns. Ensemble methods can further enhance robustness by combining multiple models.

Hyperparameter tuning should be conducted using k-fold cross-validation on the US dataset, with careful attention to ensuring the validation splits are representative. If a small subset of European data is available, it can be used for fine-tuning to improve generalization. Bayesian optimization or grid search can efficiently explore the hyperparameter space, focusing on key parameters such as regularization terms, learning rates, and model complexity. For evaluation, the primary metric should be the F1 score, which balances precision and recall and is particularly suitable for imbalanced datasets. Secondary metrics like AUC-ROC, precision, recall, and confusion matrices provide additional insights into model performance, while calibration metrics such as the Brier score ensure probabilistic predictions are reliable.

During testing, the model should be evaluated on the European dataset without further tuning, and performance metrics should be reported alongside cross-validation results from the US dataset. Error analysis is crucial to identify misclassified cases and understand whether errors stem from feature distribution differences or other factors. Additionally, the team must consider bias and fairness, ensuring the model does not disproportionately misclassify specific subgroups, and prioritize interpretability using methods like SHAP or LIME to align predictions with clinical knowledge. Continuous monitoring of the model's performance post-deployment is also recommended to detect any data drift over time. By integrating these steps, the team can develop a model that generalizes effectively to the European population while maintaining robustness, fairness, and transparency.

### 1.2 Question 2

#### 1.2.1 Part A

We need to pre-process the data. There are columns with many missing values, which should be considered to drop first. Then we can drop rows with missing values. If we simply drop rows with missing data, or drop the rows with missing data before deleting columns with significant number of missing data, it will distorted the regression results.

The result I got using self-coded  $K$ -fold CV from scratch is as below:

Best RBF Kernel Parameters:  $\lambda=0.01$ ,  $\gamma=0.1$

Best Polynomial Kernel Parameters:  $\lambda=0.01$ ,  $\text{degree}=2$

RBF Kernel Test MSE: 0.014702546829905011

Polynomial Kernel Test MSE: 0.009525269150252024

#### 1.2.2 Part B

And comparing the results of scratch code and `sklearn` code, the hyperparameters are the same:

```

Custom RBF Kernel Parameters: lambda=0.01, gamma=0.1
Custom Polynomial Kernel Parameters: lambda=0.01, degree=2
scikit-learn RBF Kernel Parameters: {'alpha': 0.01, 'gamma': 0.1}
scikit-learn Polynomial Kernel Parameters: {'alpha': 0.01, 'degree': 2}

```

The tuned parameters are the same for the built-in model of `sklearn` and the cross-validation I coded up by myself. Plugging those hyper-parameters back, we would get the same MSE as we did in Part A.

### 1.2.3 Part C

I redid the OLS model in Homework 1, and calculated the MSE using code below

```

1 residuals = y_new - y_pred
2
3 # Calculate MSE
4 mse = np.mean(residuals**2)

```

The result is as below:

```
Mean Squared Error (MSE): 0.0165
```

Here the non-linear model performed better than the linear model. First of all, linear model is a simple way to fit the data, which is highly likely to be influenced by the noise in the dataset, even though it might has lower MSE. It might cause the problem of overfitting. Whereas the CV can split the dataset into several subsets, and evaluate the model on multiple test and training sets, which would make it much generalizable, and overcome the problem of overfitting. And as we mentioned before, linear models just build the model based on data in the dataset, if the training set is not representative, it might not be a good model; but the CV is more generalizable, giving it ability to comprehensively reflect model performance.

Secondly, CV has hyper-parameters tuning, from which the optimal hyper-parameters can be selected, improving model performance.

Another aspect is that OLS lacks a regularization mechanism, making it susceptible to issues like high-dimensional data or multicollinearity. Where as we implemented regularization here, which helps control model complexity, prevent overfitting, and improve generalization. As for the efficiency to use data, the OLS just use one set of data, whereas CV divides the data multiple times, ensuring that all data points are used for both training and validation.

Finally, if we simply compare the MSE, which is the prediction errors given by different models, the MSE of cross-validated models are still smaller than those of the OLS model.

## 1.3 Question 3

### 1.3.1 Part A

Here I listed the confusion matrices, test accuracies given by python for 5 methods below. I did not generate the confusion matrix picture due to the restriction of page-count. Full reports, including `recall` rate, `f1-score`, etc, were inserted in the Appendix by me.

Firstly, for Logistic Regression in One-vs-Rest setting, I got the report

```

Test Accuracy: 93.54%
Confusion Matrix:
[[1328  39  52]
 [  54 1195  45]
 [  39  50 1254]]

```

For Multinomial Regression

```

Test Accuracy: 93.12%
Confusion Matrix:
[[1324  43  52]
 [  51 1201  42]
 [  39  52 1252]]

```

For Naive Bayes

```

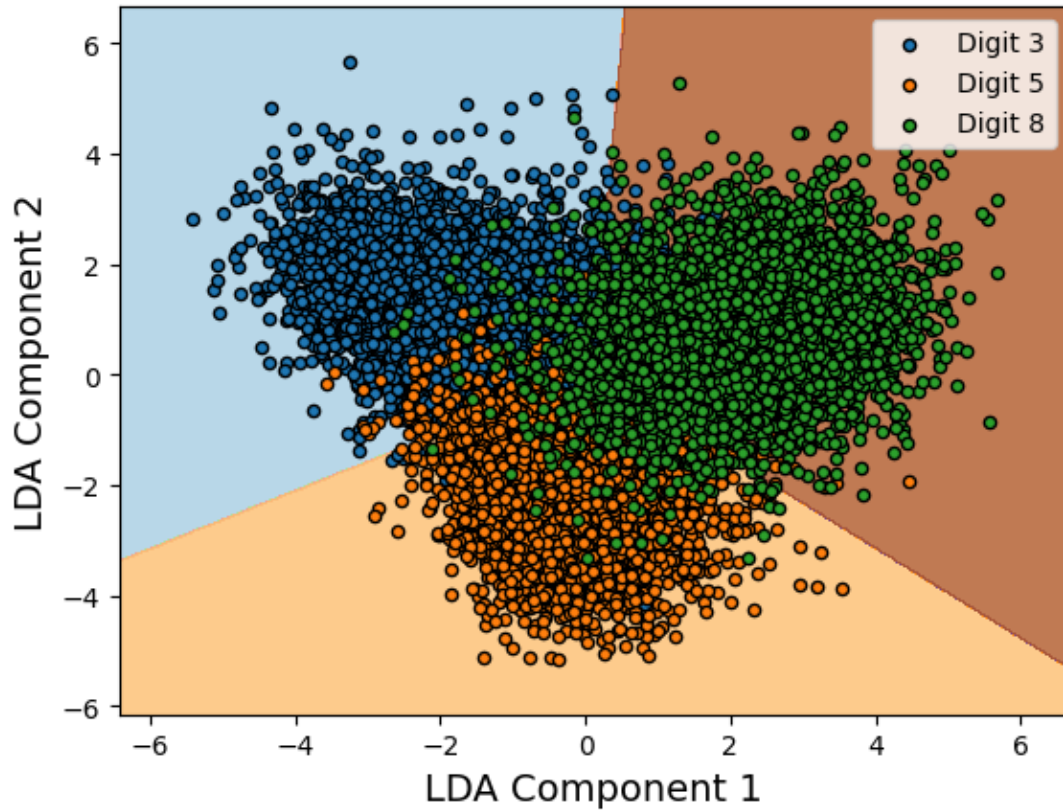
Test Accuracy: 50.22%
Confusion Matrix:
[[ 595  26 798]
 [  95 138 1061]
 [  18  21 1304]]

```

For Linear Discriminant Analysis

Test Accuracy (LDA): 91.69%  
 Confusion Matrix (LDA):  
 [[1285 69 65]  
 [ 47 1202 45]  
 [ 34 77 1232]]


Figure 1: Decision Boundary after LDA



Finally, for Linear Support Vector Machine:  
 Test Accuracy (Linear SVM): 92.73%  
 Confusion Matrix (Linear SVM):  
 [[1319 44 56]  
 [ 57 1193 44]  
 [ 39 55 1249]]

If we want to plot the decision boundary of Linear SVM, we must first do a Principle Component Analysis (PCA).

Figure 2: Decision Boundary after PCA



Overall, if we directly compare the test accuracy evaluated by python, the logistic regression has the highest accuracy rate, whereas the Naive Bayes has the least accuracy.

In the python built-in function `confusion matrix()`, the vertical axis represents the true values, whereas the horizontal axis represents the predictive value. Thus we can added the misclassified terms accordingly. However, I noticed that there was a large proportion of data misclassified in Naive Bayes case, so I splitted them into two parts, firstly I calculated the total misinterpreted values for 4 methods other than Naive Bayes; here the most often misclassified digit is 3, and has 420 counts. And for 5 and 8, the numbers are equal, which is 385.

If we add up the NB method, overall, digit 5 is more likely to be misclassified, and this trend towards error reaches to peak when we use Naive Bayes method. The total number of misclassified 5 is 1542, and 3 is 1244, and 8 is 424.

Figure 3: Count of Misclassified Digits


	A	B	C	D	E	F	G	H	I
1		<b>misclassified</b>	<b>Logistic</b>	<b>Multinomial</b>	<b>LDA</b>	<b>LSVM</b>	<b>Total, except NB</b>	<b>Naive Bayes</b>	<b>Total, with NB</b>
2		mis-classified as 5	39	43	69	44	195	26	221
3	<b>3</b>	mis-classified as 8	52	52	65	56	225	798	1023
4		total	91	95	134	100	420	824	1244
5		mis-classified as 3	54	51	47	57	209	96	305
6	<b>5</b>	mis-classified as 8	45	42	45	44	176	1061	1237
7		total	99	93	92	101	385	1157	1542
8		mis-classified as 3	39	39	34	39	151	18	169
9	<b>8</b>	mis-classified as 5	50	52	77	55	234	21	255
10		total	89	91	111	94	385	39	424
11									

### 1.3.2 Part B

The group-lasso regularized multinomial logistic regression will help us to select features that separate the three digits, and we can end up getting a heatmap. In the heatmap, bright regions represent pixels that are highly important for distinguishing between the digits 3, 5, and 8. The model relies heavily on these pixels for making predictions.

And dark regions represent pixels that contribute little to the classification task. The model considers these pixels unimportant.

Figure 4: Headmap showing Features that Separate the Three Digits



## 2 Appendix

Supplementary Github Repo

### 2.1 Question 2

#### 2.1.1 Setup

```
1      # Same as Homework 1
2
3      from ucimlrepo import fetch_ucirepo
4
5
6      # fetch dataset
7      communities_and_crime = fetch_ucirepo(id=183)
8
9      # data (as pandas dataframes)
10     X = communities_and_crime.data.features
11     y = communities_and_crime.data.targets
12
13     # metadata
14     print(communities_and_crime.metadata)
15
16     # variable information
17     print(communities_and_crime.variables)
18
19     # Inspect the shape of X and y
20     print(X.shape) # Should be (1994, 127)
21     print(y.shape) # Should be (1994, 1)
22
23     # Check for missing values
24     print(X.isnull().sum()) # Count of missing values per feature
25
26     # Inspect the first few rows of X and y
27     print(X.head())
28     print(y.head())
29     X = X.iloc[:, 5:]
30     print(X.dtypes) # There are object columns within the data. The object data type is
                     # the default type for columns containing text (strings) in a pandas DataFrame.
```

Drop columns and rows with missing values:

```
1      # Convert all values to numeric (force non-numeric to NaN)
2      X = X.applymap(pd.to_numeric, errors='coerce')
3
4      # Replace "?" with NaN
5      X.replace("?", np.nan, inplace=True)
6
7      # Check the number of missing values in each column
8      missing_counts = X.isnull().sum()
9
10     # Determine threshold for column removal (e.g., remove if >50% missing)
11     threshold = 0.5 * len(X) # Adjust this threshold as needed
12     cols_to_drop = missing_counts[missing_counts > threshold].index
13
14     # Drop those columns
15     X_cleaned = X.drop(columns=cols_to_drop)
16
17     X_with_y = pd.concat([X_cleaned, y], axis=1)
18
19     # Remove rows containing missing values in the combined DataFrame
20     X_with_y_cleaned = X_with_y.dropna()
21
```

```

22     X_with_y_cleaned
23
24     # Separate X and y after cleaning
25     X_final = X_with_y_cleaned.drop(columns=['ViolentCrimesPerPop'])
26     y_final = X_with_y_cleaned['ViolentCrimesPerPop']

```

Standardize it:

```

1     Y = y_final.to_numpy()
2
3     # Standardize X (centered and scaled)
4     scaler = StandardScaler(with_mean=True, with_std=True)
5     X_standardized = scaler.fit_transform(X_final)

```

### 2.1.2 Split into Different Sets

```

1     # Set a random seed for reproducibility
2     rng = default_rng(1)
3
4     # Define the proportions for the split
5     train_prop = 0.6
6     validation_prop = 0.2
7     test_prop = 0.2
8
9     # Calculate the number of observations for each split
10    total_samples = X_with_y_cleaned.shape[0]
11    train_size = int(train_prop * total_samples)
12    validation_size = int(validation_prop * total_samples)
13    test_size = total_samples - train_size - validation_size
14
15    # Create a random permutation of row indices
16    indices = rng.choice(np.arange(total_samples), size=(total_samples), replace=False)
17
18    # Split the dataset into train, validation, and test sets
19    y_train = Y[indices[:train_size]]
20    y_val = Y[indices[(train_size + 1):(train_size + validation_size)]]
21    y_test = Y[indices[(train_size + validation_size + 1):]]
22
23    X_train = X_standardized[indices[:train_size]]
24    X_val = X_standardized[indices[(train_size + 1):(train_size + validation_size)]]
25    X_test = X_standardized[indices[(train_size + validation_size + 1):]]

```

### 2.1.3 Using K-fold CV from scratch to tune Kernel Ridge Regularization

```

1     import numpy as np
2     from sklearn.metrics import mean_squared_error
3
4     # Define the RBF kernel function
5     def rbf_kernel(X1, X2, gamma):
6         pairwise_dists = np.sum(X1**2, axis=1).reshape(-1, 1) + np.sum(X2**2, axis=1) - 2 *
            np.dot(X1, X2.T)
7         return np.exp(-gamma * pairwise_dists)
8
9     # Define the Polynomial kernel function
10    def polynomial_kernel(X1, X2, degree, c=1):
11        return (np.dot(X1, X2.T) + c) ** degree
12
13    # Kernel Ridge Regression
14    def kernel_ridge_regression(X_train, y_train, X_test, kernel, lambda_, **
        kernel_params):
15        """

```

```

16 Perform Kernel Ridge Regression.
17
18 Parameters:
19 X_train: Training data (n_samples, n_features).
20 y_train: Target values (n_samples,).
21 X_test: Test data (n_samples_test, n_features).
22 kernel: Kernel function (rbf_kernel or polynomial_kernel).
23 lambda_: Regularization parameter.
24 **kernel_params: Kernel-specific parameters (e.g., gamma for RBF, degree for
    Polynomial).
25
26 Returns:
27 y_pred: Predicted values for X_test.
28 """
29 # Compute the kernel matrix
30 K_train = kernel(X_train, X_train, **kernel_params)
31 K_test = kernel(X_test, X_train, **kernel_params)
32
33 # Solve for the dual coefficients alpha
34 n_samples = X_train.shape[0]
35 alpha = np.linalg.inv(K_train + lambda_ * np.eye(n_samples)) @ y_train
36
37 # Predict on the test set
38 y_pred = K_test @ alpha
39 return y_pred
40
41 # K-Fold Cross-Validation
42 def k_fold_cross_validation(X_standardized, Y, k, kernel, lambda_range, gamma_range=
    None, degree_range=None):
43     """
44     Perform K-fold cross-validation to select the best hyperparameters.
45
46     Parameters:
47     X: Input data (n_samples, n_features).
48     y: Target values (n_samples,).
49     k: Number of folds.
50     kernel: Kernel function (rbf_kernel or polynomial_kernel).
51     lambda_range: List of regularization parameters to try.
52     gamma_range: List of gamma values for RBF kernel (optional).
53     degree_range: List of degree values for Polynomial kernel (optional).
54
55     Returns:
56     best_lambda: Best regularization parameter.
57     best_gamma: Best gamma value (for RBF kernel).
58     best_degree: Best degree value (for Polynomial kernel).
59     """
60     fold_size = len(X_standardized) // k
61     best_lambda = None
62     best_gamma = None
63     best_degree = None
64     best_score = float('inf')
65
66     # Grid search over hyperparameters
67     for lambda_ in lambda_range:
68         if kernel == rbf_kernel:
69             # RBF kernel: only gamma is needed
70             for gamma in (gamma_range if gamma_range is not None else [1.0]):
71                 scores = []
72                 for i in range(k):
73                     # Split into training and validation sets
74                     val_indices = range(i * fold_size, (i + 1) * fold_size)
75                     train_indices = np.setdiff1d(range(len(X)), val_indices)

```



```

77 X_train, y_train = X[train_indices], y[train_indices]
78 X_val, y_val = X[val_indices], y[val_indices]
79
80 # Train and predict
81 y_pred = kernel_ridge_regression(X_train, y_train, X_val, kernel, lambda_, gamma=
      gamma)
82
83 # Compute the validation score (e.g., mean squared error)
84 score = mean_squared_error(y_val, y_pred)
85 scores.append(score)
86
87 # Average score across folds
88 avg_score = np.mean(scores)
89 if avg_score < best_score:
90     best_score = avg_score
91     best_lambda = lambda_
92     best_gamma = gamma
93
94 elif kernel == polynomial_kernel:
95     # Polynomial kernel: degree and optionally c are needed
96     for degree in (degree_range if degree_range is not None else [2]):
97         scores = []
98         for i in range(k):
99             # Split into training and validation sets
100             val_indices = range(i * fold_size, (i + 1) * fold_size)
101             train_indices = np.setdiff1d(range(len(X)), val_indices)
102
103             X_train, y_train = X[train_indices], y[train_indices]
104             X_val, y_val = X[val_indices], y[val_indices]
105
106             # Train and predict
107             y_pred = kernel_ridge_regression(X_train, y_train, X_val, kernel, lambda_, degree=
                  degree)
108
109             # Compute the validation score (e.g., mean squared error)
110             score = mean_squared_error(y_val, y_pred)
111             scores.append(score)
112
113             # Average score across folds
114             avg_score = np.mean(scores)
115             if avg_score < best_score:
116                 best_score = avg_score
117                 best_lambda = lambda_
118                 best_degree = degree
119
120         return best_lambda, best_gamma, best_degree
121
122 # Example usage
123 if __name__ == "__main__":
124     # Generate synthetic data
125     np.random.seed(42)
126     X = np.random.rand(100, 2) # 100 samples, 2 features
127     y = 3 * X[:, 0] + 5 * X[:, 1] + np.random.randn(100) * 0.1 # Linear relationship
          with noise
128
129     # Split data into train, validation, and test sets
130     train_size = int(0.6 * len(X))
131     val_size = int(0.2 * len(X))
132     test_size = len(X) - train_size - val_size
133
134     indices = np.random.permutation(len(X))
135     X_train, y_train = X[indices[:train_size]], y[indices[:train_size]]

```

```

136 X_val, y_val = X[indices[train_size:train_size + val_size]], y[indices[train_size:
      train_size + val_size]]
137 X_test, y_test = X[indices[train_size + val_size:]], y[indices[train_size + val_size
      :]]
138
139 # Define hyperparameter ranges
140 lambda_range = [0.01, 0.1, 1, 10]
141 gamma_range = [0.01, 0.1, 1] # For RBF kernel
142 degree_range = [2, 3, 4, 5, 6, 7] # For Polynomial kernel
143
144 # Perform K-fold cross-validation for RBF kernel
145 best_lambda_rbf, best_gamma_rbf, _ = k_fold_cross_validation(
146 X_train, y_train, k=5, kernel=rbf_kernel, lambda_range=lambda_range, gamma_range=
      gamma_range
147 )
148 print(f"Best RBF Kernel Parameters: lambda={best_lambda_rbf}, gamma={best_gamma_rbf}"
      )
149
150 # Perform K-fold cross-validation for Polynomial kernel
151 best_lambda_poly, _, best_degree_poly = k_fold_cross_validation(
152 X_train, y_train, k=5, kernel=polynomial_kernel, lambda_range=lambda_range,
      degree_range=degree_range
153 )
154 print(f"Best Polynomial Kernel Parameters: lambda={best_lambda_poly}, degree={
      best_degree_poly}")
155
156 # Evaluate on the test set with the best parameters
157 y_pred_rbf = kernel_ridge_regression(X_train, y_train, X_test, rbf_kernel,
      best_lambda_rbf, gamma=best_gamma_rbf)
158 y_pred_poly = kernel_ridge_regression(X_train, y_train, X_test, polynomial_kernel,
      best_lambda_poly, degree=best_degree_poly)
159
160 print(f"RBF Kernel Test MSE: {mean_squared_error(y_test, y_pred_rbf)}")
161 print(f"Polynomial Kernel Test MSE: {mean_squared_error(y_test, y_pred_poly)}")

```

### 2.1.4 Redo the OLS as in Homework 1

```

1 # Step 1: Separate y and X
2 X_new = X_final
3 y_new = y_final
4
5 # Step 2: Add a constant to X (for the intercept term)
6 X_new = sm.add_constant(X_new)
7
8 # Step 3: Fit the OLS model
9 model = sm.OLS(y_new, X_new)
10
11 results = model.fit()
12
13 # Step 4: View the results
14 print(results.summary())

```

# OLS Regression Results

```

=====
Dep. Variable:      ViolentCrimesPerPop      R-squared:                0.696
Model:              OLS                      Adj. R-squared:           0.680
Method:             Least Squares            F-statistic:              43.26
Date:              Fri, 07 Mar 2025          Prob (F-statistic):       0.00
Time:              13:15:31                  Log-Likelihood:           1261.1
No. Observations:   1993                    AIC:                     -2320.
Df Residuals:       1892                    BIC:                     -1755.
Df Model:           100
Covariance Type:    nonrobust
=====

```

coef	std err	t	P> t	[0.025	0.975]	
const		0.5504	0.203	2.712	0.007	0.152 0.948
population		0.1840	0.397	0.463	0.643	-0.595 0.963
householdsize		-0.0223	0.086	-0.259	0.796	-0.191 0.147
racepctblack		0.2049	0.051	4.008	0.000	0.105 0.305
racePctWhite		-0.0492	0.059	-0.837	0.403	-0.164 0.066
racePctAsian		-0.0144	0.034	-0.420	0.674	-0.082 0.053
racePctHisp		0.0609	0.053	1.139	0.255	-0.044 0.166
agePct12t21		0.1104	0.106	1.043	0.297	-0.097 0.318
agePct12t29		-0.2292	0.156	-1.467	0.143	-0.536 0.077
agePct16t24		-0.1302	0.164	-0.793	0.428	-0.452 0.192
agePct65up		0.0497	0.103	0.481	0.630	-0.153 0.253
numbUrban		-0.2964	0.387	-0.766	0.444	-1.055 0.462
pctUrban		0.0467	0.016	2.989	0.003	0.016 0.077
medIncome		-0.1998	0.173	-1.158	0.247	-0.538 0.139
pctWWage		-0.2016	0.089	-2.259	0.024	-0.377 -0.027
pctWFarmSelf		0.0488	0.020	2.422	0.016	0.009 0.088
pctWInvInc		-0.1731	0.068	-2.563	0.010	-0.306 -0.041
pctWSocSec		0.0762	0.107	0.712	0.477	-0.134 0.286
pctWPubAsst		0.0050	0.046	0.108	0.914	-0.085 0.095
pctWRetire		-0.0900	0.037	-2.445	0.015	-0.162 -0.018
medFamInc		0.2880	0.160	1.797	0.073	-0.026 0.602
perCapInc		0.0955	0.189	0.506	0.613	-0.274 0.465
whitePerCap		-0.3510	0.152	-2.303	0.021	-0.650 -0.052
blackPerCap		-0.0288	0.025	-1.131	0.258	-0.079 0.021
indianPerCap		-0.0357	0.019	-1.841	0.066	-0.074 0.002
AsianPerCap		0.0216	0.019	1.145	0.252	-0.015 0.059
OtherPerCap		0.0438	0.019	2.341	0.019	0.007 0.081
HispPerCap		0.0357	0.025	1.435	0.151	-0.013 0.085
NumUnderPov		0.1112	0.138	0.805	0.421	-0.160 0.382
PctPopUnderPov		-0.1721	0.063	-2.745	0.006	-0.295 -0.049
PctLess9thGrade		-0.0999	0.068	-1.474	0.141	-0.233 0.033
PctNotHSGrad		0.0525	0.096	0.548	0.584	-0.136 0.241
PctBSorMore		0.0504	0.077	0.651	0.515	-0.101 0.202
PctUnemployed		0.0045	0.041	0.111	0.911	-0.075 0.084
PctEmploy		0.2485	0.079	3.151	0.002	0.094 0.403
PctEmplManu		-0.0658	0.032	-2.054	0.040	-0.129 -0.003
PctEmplProfServ		-0.0267	0.041	-0.654	0.513	-0.107 0.053
PctOccupManu		0.0723	0.055	1.318	0.188	-0.035 0.180
PctOccupMgmtProf		0.1226	0.086	1.419	0.156	-0.047 0.292
MalePctDivorce		0.4585	0.248	1.851	0.064	-0.027 0.944
MalePctNevMarr		0.2267	0.068	3.339	0.001	0.094 0.360
FemalePctDiv		0.1627	0.309	0.526	0.599	-0.444 0.770
TotalPctDiv		-0.5619	0.519	-1.084	0.279	-1.579 0.455
PersPerFam		-0.1405	0.168	-0.834	0.404	-0.471 0.190
PctFam2Par		0.0186	0.160	0.117	0.907	-0.294 0.331
PctKids2Par		-0.3227	0.155	-2.080	0.038	-0.627 -0.018
PctYoungKids2Par		-0.0323	0.048	-0.670	0.503	-0.127 0.062
PctTeen2Par		-0.0029	0.043	-0.069	0.945	-0.087 0.081

## 2.2 Question 3

### 2.2.1 Read the Data

```
1 import numpy as np
2 from sklearn import datasets
3 from sklearn.model_selection import train_test_split
4 from sklearn.preprocessing import StandardScaler
5 import matplotlib.pyplot as plt
6 from sklearn.decomposition import PCA
7
8 # Load MNIST dataset and filter digits 3, 5, and 8
9 from sklearn.datasets import fetch_openml
10
11 # Fetch MNIST from openml
12 mnist = fetch_openml('mnist_784', version=1, as_frame=False)
13 X, y = mnist["data"], mnist["target"]
14
15 mnist_df = pd.DataFrame(np.concatenate((mnist['target'].reshape(-1, 1), mnist['data',
16 ]), axis=1),
17 columns= ['target'] + mnist['feature_names'])
18
19 mnist_df
```

Then we keep only 3, 5, and 8.

```
1 import numpy as np
2 import pandas as pd
3 from sklearn.datasets import fetch_openml
4
5 # Fetch MNIST from OpenML
6 mnist = fetch_openml('mnist_784', version=1, as_frame=False)
7 X, y = mnist["data"], mnist["target"]
8
9 # Convert labels to integers
10 y = y.astype(int)
11
12 # Correct filtering: Keep only digits 3, 5, and 8
13 selected_digits = {3, 5, 8}
14 filter_mask = np.isin(y, list(selected_digits))
15
16 X_filtered = X[filter_mask] # Keep only selected digits
17 y_filtered = y[filter_mask] # Keep corresponding labels
18
19 # Print dataset size
20 print(f"Original dataset size: {X.shape[0]}")
21 print(f"Filtered dataset size (only 3, 5, 8): {X_filtered.shape[0]}")
22
23 # Convert to DataFrame (Optional)
24 mnist_df = pd.DataFrame(np.column_stack((y_filtered, X_filtered)),
25 columns=['target'] + mnist.feature_names)
26
27 # Display the dataframe
28 mnist_df
```

### 2.2.2 Logistic Regression with OvR

```
1 from sklearn.datasets import fetch_openml
2 from sklearn.model_selection import train_test_split
3 from sklearn.linear_model import LogisticRegression
4 from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
5 import numpy as np
```

```

6 import matplotlib.pyplot as plt
7
8 # 1. Load MNIST
9 mnist = fetch_openml('mnist_784', version=1, as_frame=False)
10 X, y = mnist.data, mnist.target
11 y = y.astype(int)
12
13 # filter 3, 5, 8
14 selected_digits = [3, 5, 8]
15 mask = np.isin(y, selected_digits)
16 X_filtered, y_filtered = X[mask], y[mask]
17
18 X_filtered = X_filtered / 255.0
19
20 # Divide the sets
21 X_train, X_test, y_train, y_test = train_test_split(X_filtered, y_filtered, test_size
    =0.2, random_state=42)
22
23 # 2. Train the Logistic Regression (One-vs-Rest)
24 log_reg_ovr = LogisticRegression(multi_class='ovr', solver='liblinear', max_iter=100,
    random_state=42)
25 log_reg_ovr.fit(X_train, y_train)
26
27 # 3. Evaluate
28 y_pred = log_reg_ovr.predict(X_test)
29 accuracy = accuracy_score(y_test, y_pred)
30 print(f"Test Accuracy: {accuracy*100:.2f}%")
31
32 # Confusion Matrix
33 conf_matrix = confusion_matrix(y_test, y_pred)
34 print("Confusion Matrix:")
35 print(conf_matrix)
36
37 # Class Report
38 class_report = classification_report(y_test, y_pred)
39 print("Classification Report:")
40 print(class_report)
41
42 # 4. Visualize
43 num_images = 5
44 indices = np.random.choice(len(X_test), num_images, replace=False)
45
46 plt.figure(figsize=(10, 5))
47 for i, index in enumerate(indices):
48     plt.subplot(1, num_images, i + 1)
49     plt.imshow(X_test[index].reshape(28, 28), cmap='gray')
50     plt.title(f"Pred: {y_pred[index]}\nTrue: {y_test[index]}")
51     plt.axis('off')
52 plt.show()

```

Test Accuracy: 93.12%

Confusion Matrix:

```
[[1328  39  52]
 [  54 1195  45]
 [  39  50 1254]]
```

Classification Report:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

3	0.93	0.94	0.94	1419
5	0.93	0.92	0.93	1294
8	0.93	0.93	0.93	1343

accuracy			0.93	4056
macro avg	0.93	0.93	0.93	4056
weighted avg	0.93	0.93	0.93	4056

### 2.2.3 Multinomial Regression

```
1 from sklearn.datasets import fetch_openml
2 from sklearn.model_selection import train_test_split
3 from sklearn.linear_model import LogisticRegression
4 from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
5 import numpy as np
6 import matplotlib.pyplot as plt
7
8 # Similar steps
9
10 mnist = fetch_openml('mnist_784', version=1, as_frame=False)
11 X, y = mnist.data, mnist.target
12 y = y.astype(int)
13
14 selected_digits = [3, 5, 8]
15 mask = np.isin(y, selected_digits)
16 X_filtered, y_filtered = X[mask], y[mask]
17
18 X_filtered = X_filtered / 255.0
19
20 X_train, X_test, y_train, y_test = train_test_split(X_filtered, y_filtered, test_size
    =0.2, random_state=42)
21
22 log_reg_multinomial = LogisticRegression(multi_class='multinomial', solver='lbfgs',
    max_iter=100, random_state=42)
23 log_reg_multinomial.fit(X_train, y_train)
24
25 y_pred = log_reg_multinomial.predict(X_test)
26 accuracy = accuracy_score(y_test, y_pred)
27 print(f"Test Accuracy: {accuracy*100:.2f}%")
28
29 conf_matrix = confusion_matrix(y_test, y_pred)
30 print("Confusion Matrix:")
31 print(conf_matrix)
32
33 class_report = classification_report(y_test, y_pred)
34 print("Classification Report:")
35 print(class_report)
36
37 num_images = 5
38 indices = np.random.choice(len(X_test), num_images, replace=False)
39
40 plt.figure(figsize=(10, 5))
41 for i, index in enumerate(indices):
42     plt.subplot(1, num_images, i + 1)
```

```

43 plt.imshow(X_test[index].reshape(28, 28), cmap='gray')
44 plt.title(f"Pred: {y_pred[index]}\nTrue: {y_test[index]}")
45 plt.axis('off')
46 plt.show()

```

Test Accuracy: 93.12%

Confusion Matrix:

```

[[1324  43  52]
 [ 51 1201  42]
 [ 39  52 1252]]

```

Classification Report:

	precision	recall	f1-score	support
3	0.94	0.93	0.93	1419
5	0.93	0.93	0.93	1294
8	0.93	0.93	0.93	1343

3	0.94	0.93	0.93	1419
5	0.93	0.93	0.93	1294
8	0.93	0.93	0.93	1343

accuracy			0.93	4056
----------	--	--	------	------

macro avg	0.93	0.93	0.93	4056
-----------	------	------	------	------

weighted avg	0.93	0.93	0.93	4056
--------------	------	------	------	------

Do the Learning Curve

```

1  import numpy as np
2  import matplotlib.pyplot as plt
3  from sklearn.datasets import fetch_openml
4  from sklearn.model_selection import train_test_split, learning_curve
5  from sklearn.linear_model import LogisticRegression
6
7  # 1. Load and filter the MNIST dataset
8  mnist = fetch_openml('mnist_784', version=1, as_frame=False)
9  X, y = mnist.data, mnist.target
10 y = y.astype(int)
11
12 # Filter out samples with labels 3, 5, 8
13 selected_digits = [3, 5, 8]
14 mask = np.isin(y, selected_digits)
15 X_filtered, y_filtered = X[mask], y[mask]
16
17 # Scale pixel values from [0, 255] to [0, 1]
18 X_filtered = X_filtered / 255.0
19
20 # Split the dataset into training and testing sets
21 X_train, X_test, y_train, y_test = train_test_split(X_filtered, y_filtered, test_size
    =0.2, random_state=42)
22
23 # 2. Define the Logistic Regression model
24 log_reg = LogisticRegression(solver='lbfgs', max_iter=100, random_state=42)
25
26 # 3. Compute the learning curve
27 train_sizes, train_scores, val_scores = learning_curve(
28 log_reg, X_train, y_train, cv=5, scoring='accuracy', train_sizes=np.linspace(0.1,
    1.0, 10)
29 )
30
31 # Calculate mean and standard deviation of training and validation scores
32 train_scores_mean = np.mean(train_scores, axis=1)
33 train_scores_std = np.std(train_scores, axis=1)
34 val_scores_mean = np.mean(val_scores, axis=1)
35 val_scores_std = np.std(val_scores, axis=1)
36
37 # 4. Plot the learning curve
38 plt.figure(figsize=(10, 6))

```

```

39 plt.plot(train_sizes, train_scores_mean, label='Training Accuracy', color='blue',
40          marker='o')
41 plt.fill_between(train_sizes, train_scores_mean - train_scores_std, train_scores_mean
42                  + train_scores_std, alpha=0.15, color='blue')
43 plt.plot(train_sizes, val_scores_mean, label='Validation Accuracy', color='green',
44          marker='o')
45 plt.fill_between(train_sizes, val_scores_mean - val_scores_std, val_scores_mean +
46                  val_scores_std, alpha=0.15, color='green')
47
48 # plt.title('Learning Curve for Logistic Regression', fontsize=16)
49 plt.xlabel('Training Set Size', fontsize=14)
50 plt.ylabel('Accuracy', fontsize=14)
51 plt.legend(loc='best')
52 plt.grid(True)
53 plt.show()

```

Do the Validation Curve

```

1  from sklearn.model_selection import validation_curve
2
3  # Define the range of hyperparameter C (inverse of regularization strength)
4  param_range = np.logspace(-4, 4, 10)
5
6  # Compute validation curve
7  train_scores, val_scores = validation_curve(
8      log_reg, X_train, y_train, param_name='C', param_range=param_range, cv=5, scoring='
9      accuracy'
10 )
11
12 # Calculate mean and standard deviation of training and validation scores
13 train_scores_mean = np.mean(train_scores, axis=1)
14 train_scores_std = np.std(train_scores, axis=1)
15 val_scores_mean = np.mean(val_scores, axis=1)
16 val_scores_std = np.std(val_scores, axis=1)
17
18 # Plot the validation curve
19 plt.figure(figsize=(10, 6))
20 plt.semilogx(param_range, train_scores_mean, label='Training Accuracy', color='blue',
21             marker='o')
22 plt.fill_between(param_range, train_scores_mean - train_scores_std, train_scores_mean
23                 + train_scores_std, alpha=0.15, color='blue')
24 plt.semilogx(param_range, val_scores_mean, label='Validation Accuracy', color='green',
25             marker='o')
26 plt.fill_between(param_range, val_scores_mean - val_scores_std, val_scores_mean +
27                 val_scores_std, alpha=0.15, color='green')
28
29 # plt.title('Validation Curve for Logistic Regression', fontsize=16)
30 plt.xlabel('Regularization Strength (C)', fontsize=14)
31 plt.ylabel('Accuracy', fontsize=14)
32 plt.legend(loc='best')
33 plt.grid(True)
34 plt.show()

```

## 2.2.4 Naive Bayes

```

1  from sklearn.datasets import fetch_openml
2  from sklearn.model_selection import train_test_split
3  from sklearn.naive_bayes import GaussianNB
4  from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
5  import numpy as np
6  import matplotlib.pyplot as plt
7

```



```

8 # 1. Load and filter the MNIST dataset
9 mnist = fetch_openml('mnist_784', version=1, as_frame=False)
10 X, y = mnist.data, mnist.target
11 y = y.astype(int)
12
13 # Filter out samples with labels 3, 5, 8
14 selected_digits = [3, 5, 8]
15 mask = np.isin(y, selected_digits)
16 X_filtered, y_filtered = X[mask], y[mask]
17
18 # Scale pixel values from [0, 255] to [0, 1]
19 X_filtered = X_filtered / 255.0
20
21 # Split the dataset into training and testing sets
22 X_train, X_test, y_train, y_test = train_test_split(X_filtered, y_filtered, test_size
    =0.2, random_state=42)
23
24 # 2. Train a Gaussian Naive Bayes model
25 naive_bayes = GaussianNB()
26 naive_bayes.fit(X_train, y_train)
27
28 # 3. Evaluate the model
29 y_pred = naive_bayes.predict(X_test)
30 accuracy = accuracy_score(y_test, y_pred)
31 print(f"Test Accuracy: {accuracy*100:.2f}%")
32
33 # Confusion matrix
34 conf_matrix = confusion_matrix(y_test, y_pred)
35 print("Confusion Matrix:")
36 print(conf_matrix)
37
38 # Classification report
39 class_report = classification_report(y_test, y_pred)
40 print("Classification Report:")
41 print(class_report)
42
43 # 4. Visualize the results
44 num_images = 5
45 indices = np.random.choice(len(X_test), num_images, replace=False)
46
47 plt.figure(figsize=(10, 5))
48 for i, index in enumerate(indices):
49     plt.subplot(1, num_images, i + 1)
50     plt.imshow(X_test[index].reshape(28, 28), cmap='gray')
51     plt.title(f"Pred: {y_pred[index]}\nTrue: {y_test[index]}")
52     plt.axis('off')
53 plt.show()

```

Test Accuracy: 50.22%

Confusion Matrix:

```
[[ 595   26  798]
 [  95  138 1061]
 [   18   21 1304]]
```

Classification Report:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

3	0.84	0.42	0.56	1419
5	0.75	0.11	0.19	1294
8	0.41	0.97	0.58	1343

accuracy			0.50	4056
macro avg	0.67	0.50	0.44	4056
weighted avg	0.67	0.50	0.45	4056

### 2.2.5 Linear Discriminant Analysis

```
1  from sklearn.datasets import fetch_openml
2  from sklearn.model_selection import train_test_split
3  from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
4  from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
5  import numpy as np
6  import matplotlib.pyplot as plt
7
8  # 1. Load and filter the MNIST dataset
9  mnist = fetch_openml('mnist_784', version=1, as_frame=False)
10 X, y = mnist.data, mnist.target
11 y = y.astype(int)
12
13 # Filter out samples with labels 3, 5, 8
14 selected_digits = [3, 5, 8]
15 mask = np.isin(y, selected_digits)
16 X_filtered, y_filtered = X[mask], y[mask]
17
18 # Scale pixel values from [0, 255] to [0, 1]
19 X_filtered = X_filtered / 255.0
20
21 # Split the dataset into training and testing sets
22 X_train, X_test, y_train, y_test = train_test_split(X_filtered, y_filtered, test_size
    =0.2, random_state=42)
23
24 # 2. Train a Linear Discriminant Analysis (LDA) model
25 lda = LinearDiscriminantAnalysis()
26 lda.fit(X_train, y_train)
27
28 # 3. Evaluate the model
29 y_pred = lda.predict(X_test)
30 accuracy = accuracy_score(y_test, y_pred)
31 print(f"Test Accuracy (LDA): {accuracy*100:.2f}%")
32
33 # Confusion matrix
34 conf_matrix = confusion_matrix(y_test, y_pred)
35 print("Confusion Matrix (LDA):")
36 print(conf_matrix)
37
38 # Classification report
39 class_report = classification_report(y_test, y_pred)
40 print("Classification Report (LDA):")
41 print(class_report)
42
43 # 4. Visualize the results
```

```

44 num_images = 5
45 indices = np.random.choice(len(X_test), num_images, replace=False)
46
47 plt.figure(figsize=(10, 5))
48 for i, index in enumerate(indices):
49     plt.subplot(1, num_images, i + 1)
50     plt.imshow(X_test[index].reshape(28, 28), cmap='gray')
51     plt.title(f"Pred: {y_pred[index]}\nTrue: {y_test[index]}")
52     plt.axis('off')
53 plt.show()

```

Test Accuracy (LDA): 91.69%

Confusion Matrix (LDA):

```

[[1285  69  65]
 [ 47 1202  45]
 [ 34  77 1232]]

```

Classification Report (LDA):

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

3	0.94	0.91	0.92	1419
5	0.89	0.93	0.91	1294
8	0.92	0.92	0.92	1343

accuracy			0.92	4056
macro avg	0.92	0.92	0.92	4056
weighted avg	0.92	0.92	0.92	4056

## 2.2.6 Linear SVM

```

1 from sklearn.datasets import fetch_openml
2 from sklearn.model_selection import train_test_split
3 from sklearn.svm import LinearSVC
4 from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
5 import numpy as np
6 import matplotlib.pyplot as plt
7
8 # 1. Load and filter the MNIST dataset
9 mnist = fetch_openml('mnist_784', version=1, as_frame=False)
10 X, y = mnist.data, mnist.target
11 y = y.astype(int)
12
13 # Filter out samples with labels 3, 5, 8
14 selected_digits = [3, 5, 8]
15 mask = np.isin(y, selected_digits)
16 X_filtered, y_filtered = X[mask], y[mask]
17
18 # Scale pixel values from [0, 255] to [0, 1]
19 X_filtered = X_filtered / 255.0
20
21 # Split the dataset into training and testing sets
22 X_train, X_test, y_train, y_test = train_test_split(X_filtered, y_filtered, test_size=0.2, random_state=42)
23
24 # 2. Train a Linear SVM model with One-vs-Rest strategy
25 linear_svm = LinearSVC(multi_class='ovr', max_iter=10000, random_state=42) # Use One-vs-Rest
26 linear_svm.fit(X_train, y_train)
27
28 # 3. Evaluate the model
29 y_pred = linear_svm.predict(X_test)
30 accuracy = accuracy_score(y_test, y_pred)

```

```

31 print(f"Test Accuracy (Linear SVM): {accuracy*100:.2f}%")
32
33 # Confusion matrix
34 conf_matrix = confusion_matrix(y_test, y_pred)
35 print("Confusion Matrix (Linear SVM):")
36 print(conf_matrix)
37
38 # Classification report
39 class_report = classification_report(y_test, y_pred)
40 print("Classification Report (Linear SVM):")
41 print(class_report)
42
43 # 4. Visualize the results
44 num_images = 5
45 indices = np.random.choice(len(X_test), num_images, replace=False)
46
47 plt.figure(figsize=(10, 5))
48 for i, index in enumerate(indices):
49     plt.subplot(1, num_images, i + 1)
50     plt.imshow(X_test[index].reshape(28, 28), cmap='gray')
51     # plt.title(f"Pred: {y_pred[index]}\nTrue: {y_test[index]}")
52     plt.axis('off')
53 plt.show()

```

Test Accuracy (Linear SVM): 92.73%

Confusion Matrix (Linear SVM):

```

[[1319  44  56]
 [ 57 1193  44]
 [ 39  55 1249]]

```

Classification Report (Linear SVM):

	precision	recall	f1-score	support
3	0.93	0.93	0.93	1419
5	0.92	0.92	0.92	1294
8	0.93	0.93	0.93	1343
accuracy			0.93	4056
macro avg	0.93	0.93	0.93	4056
weighted avg	0.93	0.93	0.93	4056

## 2.2.7 Plot the Confusion Matrix

```

1 from sklearn.metrics import confusion_matrix
2 import seaborn as sns
3
4 # Logistic Regression Confusion Matrix
5 log_reg_pred = log_reg.predict(X_test)
6 log_reg_cm = confusion_matrix(y_test, log_reg_pred)
7
8 # Multinomial Regression Confusion Matrix
9 multinomial_pred = multinomial_reg.predict(X_test)
10 multinomial_cm = confusion_matrix(y_test, multinomial_pred)
11
12 # Naive Bayes Confusion Matrix
13 naive_bayes_pred = naive_bayes.predict(X_test)
14 naive_bayes_cm = confusion_matrix(y_test, naive_bayes_pred)
15
16 # LDA Confusion Matrix
17 lda_pred = lda.predict(X_test)
18 lda_cm = confusion_matrix(y_test, lda_pred)
19

```

```

20 # SVM Confusion Matrix
21 svm_pred = svm.predict(X_test)
22 svm_cm = confusion_matrix(y_test, svm_pred)
23
24 def plot_confusion_matrix(cm, class_names):
25     plt.figure(figsize=(6, 5))
26     sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=class_names,
27                 yticklabels=class_names)
28     plt.xlabel('Predicted')
29     plt.ylabel('True')
30     plt.title('Confusion_Matrix')
31     plt.show()
32
33     class_names = ['3', '5', '8']
34
35 # Plot confusion matrices
36 plot_confusion_matrix(log_reg_cm, class_names)
37 plot_confusion_matrix(multinomial_cm, class_names)
38 plot_confusion_matrix(naive_bayes_cm, class_names)
39 plot_confusion_matrix(lda_cm, class_names)
40 plot_confusion_matrix(svm_cm, class_names)

```